

Table of Contents

- 1.1 Project Objectives ← covers Creativity & Innovation (7 marks)
 - 1.1.1 Objective 1 — Matrix Engine
 - 1.1.2 Objective 2 — OCC Layer
 - 1.1.3 Objective 3 — FCFS Dashboard
 - 1.1.4 Objective 4 — RBAC
 - 1.1.5 Objective 5 — Deployment
 - 1.2 Problem Statement ← covers Problem & Scope (7 marks)
(5 pain points + citations)
 - 1.3 Project Background ← context for the problem
(TAR UMT Sabah, FOCS, current workflow, venue dataset)
 - 1.4 Advantages and Contributions ← covers Contribution & Reach (6 marks)
 - 1.4.1 Advantages (over Google Sheets)
 - 1.4.2 Contributions (stakeholders, SDG 4 & 8, commercialisation)
 - 1.5 Project Plan ← covers Feasibility (5 marks)
 - 1.5.1 Functional Modules ← scope as system decomposition (like your senior)
 - 1.5.2 Milestones (Gantt table)
 - 1.5.3 Software Development Model (Agile)
 - 1.6 Chapter Summary and Evaluation
-

1.1 Project Objectives

← covers Creativity & Innovation (7 marks)

- 1.1.1 Objective 1 — Matrix Engine
- 1.1.2 Objective 2 — OCC Layer
- 1.1.3 Objective 3 — FCFS Dashboard
- 1.1.4 Objective 4 — RBAC
- 1.1.5 Objective 5 — Deployment

1.1 Project Objectives

Opening paragraph (1–2 sentences):

This section defines five project objectives formulated using the SMART framework — Specific, Measurable, Achievable, Relevant, and Time-bound within the FYP1–FYP2 timeline. Each objective is linked to a distinct technical contribution that distinguishes this system from generic scheduling solutions.

1.1.1 Objective 1 — Multi-Entity Matrix Intersection Engine

Statement:

To design and implement a Multi-Entity Matrix Intersection Engine that computes valid replacement windows by performing a four-way set intersection across lecturer availability, stacked cohort free schedules (supporting mixed-cohort groupings), room occupancy vectors, and capacity-aware venue filtering that excludes rooms with insufficient seating.

Why it scores Creativity & Innovation (7 marks):

- **New/unique:** No existing system at TAR UMT Sabah performs automated 4-vector intersection for class replacement. Spreadsheets require manual overlay.
- **Fresh approach:** Uses deterministic set intersection (not heuristic or ML-based), making results predictable and auditable.
- **Leverages IT trends:** Relational set operations on real-time timetable data, a modern alternative to static cross-referencing.

Measurable success criteria:

- Engine returns correct common free slots for any combination of multiple RSD3 cohorts within 500ms
 - Capacity filter correctly excludes rooms smaller than total enrolled student count
 - **Handles edge cases:**
 - **No common slot found** — a lecturer wants a replacement but the engine finds zero overlapping free time across lecturer, cohorts, and room. The system should return "No available slots" rather than crashing or returning incorrect results.
 - **Single-cohort request** — a replacement affects only one cohort (not mixed). The engine should still work correctly with just 3 vectors (lecturer $\times$ 1 cohort $\times$ room) instead of 4.
 - **All-day occupancy** — the lecturer is fully booked the entire day with no free slots. The engine should correctly return empty results rather than showing occupied slots as available.
-

1.1.2 Objective 2 — Optimistic Concurrency Control (OCC) Layer

Statement:

To implement an Optimistic Concurrency Control layer that prevents double-booking by performing a millisecond-precision transactional validation at submission time, automatically aborting and rolling back if a concurrent transaction has modified the slot state.

Why it scores Creativity & Innovation:

- **New/unique:** First application of OCC (rather than pessimistic locking) for class replacement at the faculty. Spreadsheets offer zero transaction support.

- **Fresh approach:** OCC avoids performance-heavy database locks, keeping the UI responsive even under concurrent lecturer submissions.
- Leverages IT trends: ACID-compliant transactional design in a web application context, referencing Kung & Robinson's foundational work.
- "The OCC layer operates within Atomicity, Consistency, Isolation, Durability (ACID) database transactions to guarantee data integrity under concurrent access."

Measurable success criteria:

- Two simultaneous submissions for the same slot → exactly one succeeds, one receives rollback alert
 - Validation checks slot state at the millisecond of database write
 - All transactional outcomes logged in audit trail
-

1.1.3 Objective 3 — FCFS Digital Approval Dashboard

Statement:

To develop a First-Come-First-Served digital approval dashboard that automates the Programme Leader's verification workflow — replacing the manual rechecking of schedules, venue cross-referencing, and "Y" annotation process — with a pre-validated queue that enables one-click approval (slot transitions to Occupied) or rejection with mandatory reason (slot returns to Available), with every action recorded in a timestamped audit trail.

Why it scores Creativity & Innovation:

- **New/unique:** Currently PLs must manually recheck every request against lecturer schedules, cohort timetables, and venue in PDFs before writing "Y" in a spreadsheet cell. The dashboard eliminates this rechecking entirely because the engine has already validated the slot.
- **Fresh approach:** One-click approval replaces a multi-step manual verification pipeline. The audit trail is automatic rather than retroactively reconstructed.
- **Leverages IT trends:** State machine design (Available → Pending → Occupied/Rejected) for structured workflow management.

Measurable success criteria:

- Queue displays all pending requests sorted by submission timestamp (earliest first — matching current chronological order)
- Each request shows pre-computed slot validity so PL does not need to re-verify
- Approval: 1 click → slot becomes Occupied, all future requests for that slot auto-blocked
- Rejection: requires mandatory text reason → slot returns to Available, available for rebooking
- Audit trail records: PL identity, timestamp, action taken, slot ID, reason (if rejected)

1.1.4 Objective 4 — Role-Based Access Control (RBAC)

Statement: To build a role-based access control system with three tiers:

- view-only access for students,
- draft-and-submit workflows for lecturers, and
- hybrid administrative privileges for Programme Leaders who retain full lecturer creation rights in addition to queue management.

Why it scores Creativity & Innovation:

- **New/unique:** Hybrid PL role (admin + lecturer creation rights) is a practical design not commonly found in generic scheduling tools. Existing Google Sheets have no role differentiation — anyone with the link can edit. The hybrid PL role (admin + lecturer rights) reflects the real FOCS structure where PLs also teach.
- **Fresh approach:** Designed around actual FOCS workflow rather than generic RBAC templates. — PLs are also teaching staff, so they need both privilege sets.
- **Leverages IT trends:** Laravel's built-in authorization gates and middleware for declarative permission checks.

Measurable success criteria:

- Student: can view timetables and replacement status only (no create/edit/delete)
- Lecturer: can create drafts, submit requests, edit own drafts. Cannot approve or view other lecturers' drafts
- PL: inherits all lecturer rights + can approve/reject any request in queue
- Authentication via username/password with session management
- Unauthenticated users: cannot access any dashboard page (redirected to login)

1.1.5 Objective 5 — Prototype Deployment

Statement:

To deploy a fully functional Laravel-PostgreSQL web application on a local development server, seeded with five RSD3 cohorts, five FOCS staff members (3 lecturers + 2 Programme Leaders), and 14 classroom venues (B005, B006, B009–B011, B101–B111) for prototype validation within the FYP2 sprint phase.

Why it scores Feasibility (5 marks) + supports Creativity:

- **Developable within FYP timeframe:** 5 cohorts, 5 staff, 14 rooms — realistic scope for 7-week sprint
- **Relevant to the problem:** Demonstrates the engine working with real-world scale data, not toy data

- **Technology:** Laravel (PHP), PostgreSQL, TailwindCSS, Blade templating — modern full-stack stack

Measurable success criteria:

- All seed data loads correctly (no missing foreign keys)
 - Application runs on local development server (Laravel Artisan serve)
 - All three user roles can log in with test credentials and access their respective dashboards
 - Each objective's success criteria individually verifiable through test cases
-

1.2 Problem Statement

Opening paragraph:

This section identifies five interrelated weaknesses in the current class replacement process at FOCS Sabah. Each problem is supported by academic literature and linked to specific consequences that motivate the development of the proposed system.

1.2.1 Fragmented Process and Human Error

Statement:

The management of replacement requests relies on a decentralised web of Google Sheets where each lecturer maintains their own schedule independently. No centralized database exists to unify the state of replacements across the faculty, requiring numerous manual cross-referencing steps for each request.

Why it reflects a real-world problem:

- Each lecturer keeps a personal schedule spreadsheet with no shared source of truth
- No central database to track replacement state across the faculty
- Every replacement request requires manual cross-referencing of multiple spreadsheets

Supporting citation:

Babaei, Karimpour, & Hadidi (2015) established that manual scheduling processes in educational organisations significantly increase the probability of human error.

Consequence: Scheduling mismatches that require downstream correction, disrupting both lecturers and students.

1.2.2 Blind Spots in Multi-Cohort Scheduling

Statement:

When a module draws students from multiple RSD3 groups simultaneously (e.g., BMSE3153 shared by G1–G5), manually combining two or more independent cohort schedules to find common free periods is both time-consuming and error-prone.

Why it reflects a real-world problem:

- Each cohort follows a distinct module timetable
- Common free period intersection is not immediately visible
- Manual overlay of separate timetable files introduces mistakes
- Errors are only discovered after replacement has been scheduled

Consequence:

Incorrect replacement slots that require further rescheduling and cause disruption to teaching activities.

1.2.3 Venue Availability Tracking

Statement:

Information regarding venue availability is stored separately from the scheduling process as a static PDF file. Lecturers must independently check whether classrooms are free for replacement, and there is no automated capacity-aware filtering. **Why it reflects a real-world problem:**

- Room availability tracked in isolation from scheduling logic
- No capacity-aware filtering to prevent assigning a room too small for combined cohorts
- Room usage priority rules (Lab only for P, B006 priority for Networking) not enforced

Consequence:

Venue conflicts discovered late, requiring additional rescheduling and administrative overhead.

1.2.4 Data Race Conditions

Statement:

The collaborative nature of shared Google Sheets offers no transaction boundaries. When multiple lecturers submit requests simultaneously, both see the same slot as available and proceed, creating an irresolvable double-booking conflict.

Why it reflects a real-world problem:

- Google Sheets has no transaction isolation between concurrent users
- Two lecturers submitting at the same time both see the slot as available
- Double-booking only discovered after both requests are processed
- No rollback mechanism exists to recover from the conflict

Supporting citations:

Al-Hawari, Al-Ashi, Abawi, & Alounch (2020) — absence of concurrency control in multi-user scheduling leads to data integrity violations. Kung & Robinson (1981) — optimistic concurrency control methods can resolve such conflicts without the performance overhead of pessimistic locking.

Consequence: Irresolvable booking conflicts requiring Programme Leader intervention and manual reconciliation.

1.2.5 Elevated Administrative Burden on Programme Leaders

Statement:

Programme Leaders must manually verify every incoming request against lecturer schedules, cohort timetables, and venue availability records before annotating approval in the spreadsheet. Each verification is performed from scratch with no automated decision support.

Why it reflects a real-world problem:

- PL manually checks each request against lecturer schedules
- PL manually cross-references cohort timetables (often multiple cohorts)
- PL manually checks venue availability from separate PDF
- All verification is performed with no automated decision support

Supporting citation:

Schaerf (1999) — centralising scheduling logic into a dedicated system eliminates the redundancy and error propagation inherent in manual coordination.

Consequence: Replacement processing takes hours or days instead of minutes, and the manual verification process itself introduces further opportunity for oversight.

Closing paragraph:

The absence of a centralised, transaction-aware scheduling platform at FOCS Sabah has resulted in recurring scheduling conflicts, prolonged replacement processing times, and an opaque audit trail, necessitating the development of an automated replacement scheduling system.

1.3 Project Background

- **TAR UMT Sabah — FOCS**
 - Faculty of Computing and Information Technology
 - Cohort-based programmes: DFT, DSF, RSD streams
 - Some modules enrol multiple cohorts together → complex scheduling interdependencies
- **Academic Staff**
 - Total: 11 lecturers (including 2 Programme Leaders who also teach)
 - Prototype dataset: **5 staff** selected from the 11
 - 3 ordinary lecturers
 - 2 Programme Lecturers (inherit full lecturer + admin privileges)
- **Cohort Structure (Full at Sabah)**
 - DFT1(S1), DFT2(S1)
 - DSF1(S1), DSF2(S1)
 - RSD1(S1)G1
 - RSD2(S1)G2, RSD2(S1)G3
 - RSD3(S1)G1 to G5
- **Prototype Dataset Selection — 5 Cohorts**
 - All 5 from RSD3(S1) G1, G2, G3, G4, G5
 - Rationale: they share core modules (BMSE3153, BMIT3173, BMIT2073, BMIS2113, BMIT3273), creating realistic multi-cohort scenarios
- **Timetable Data**
 - Source: past semester aSc Timetables export from FOCS Sabah
 - Period: 7-week short semester
 - Schedule: **identical timetable structure each week** (excludes replacement classes)
- **Physical Venues**
 - Tutorial rooms: B014–B018, B101–B109
 - Lecture halls: B110, B111
 - Computer labs: B009–B011, B005, B006 (Cisco Lab)

- All rooms found in the reference PDFs match this set
- **Physical Venues (Dataset — 14 rooms)**
 - Tutorial Rooms (capacity ≤ 35): B101, B102, B103, B104, B105, B106, B107, B108, B109
 - Lecture Halls (capacity > 35): B110, B111
 - Computer Labs (capacity 28, P only): B005, B009, B010, B011
 - Cisco Lab (capacity 32, L/T/P for Networking, P for IoT): B006
 - Excluded from dataset: B001–B004, B007–B008, B012–B018
- **Current Replacement Workflow (Manual)**
 - Entirely decentralised via shared Google Sheets
 - Steps:
 1. Lecturer consults personal timetable
 2. Cross-references schedules of all affected cohorts to find common free slots
 3. Checks venue availability from a separate PDF file
 4. Waits for PL validation and approval
 - Multi-cohort → manual overlay of separate files → highly error-prone
- **Why the Current Process is Insufficient**
 - No centralised database → each lecturer keeps their own replacement records
 - No automated cross-cohort intersection → free periods calculated by hand
 - Room availability tracked in isolation from scheduling logic
 - No transaction boundaries → concurrent double-booking possible
 - PLs manually verify every request against lecturer schedules, cohort timetables, and venue records
 - Processing a single replacement is very time-consuming (hours to days)

Project Team and Organization

This project is undertaken by Poong Foo Jing (student) under the supervision of Mr. Lim Jia Zheng, with Ms. Teng Nga Sing as moderator. Domain expertise and requirements validation are provided by two FOCS Programme Leaders: Pn. Surayaini Binti Basri and En. Mohd Nur Rahmat Bin Mohd Taat. The client is the Faculty of Computing and Information Technology, TAR UMT Sabah.

1.4 Advantages and Contributions

← covers Contribution & Reach (6 marks)

1.4.1 Advantages (over Google Sheets)

1.4.2 Contributions (stakeholders, SDG 4 & 8, commercialisation)

Opening paragraph:

This section outlines the advantages of the proposed TARUMT Class Replacement System over the existing Google Sheets-based workflow, followed by its contributions to stakeholders, the institution, and alignment with the United Nations Sustainable Development Goals.

1.4.1 Advantages (over Google Sheets)

Advantage 1 — Eliminates Manual Cross-Referencing

- Google Sheets: lecturers manually overlay separate timetable files to find common free periods
- Proposed system: Matrix Intersection Engine computes valid slots automatically in milliseconds
- Benefit: reduces replacement processing from hours/days to minutes

Advantage 2 — Prevents Double-Booking

- Google Sheets: no transaction boundaries → simultaneous submissions cause conflicts
- Proposed system: OCC validates each submission at database-write level → double-booking impossible
- Benefit: eliminates administrative overhead of resolving booking conflicts

Advantage 3 — Automated Venue Verification

- Google Sheets: lecturers check venue availability from separate static PDF
- Proposed system: capacity-aware room filtering is integrated into the intersection engine
- Benefit: venues with insufficient seating are automatically excluded

Advantage 4 — Centralised Replacement State

Google Sheets (current):

- Each lecturer maintains their own personal schedule Google Sheet
- There is no central record of which replacement requests exist, who submitted them, or their current status
- If Lecturer A submits a request, Lecturer B has no way of knowing unless the PL manually communicates it
- The state of a slot (free, pending, approved, occupied by another booking) is not tracked anywhere
- At any given time, no single person has a complete picture of all replacements happening in the faculty

Proposed system:

- All timetable data, replacement requests, and slot states live in a single PostgreSQL

database

- Every slot has an explicit state:

State	Meaning
Available	Slot is free — no conflict exists, no replacement pending
Pending (Self)	The logged-in lecturer has submitted a request for this slot, awaiting PL approval
Reserved (Other)	Another lecturer has submitted a request for this slot (FCFS priority applies)
Occupied	PL has approved the replacement — slot is locked permanently

Benefit:

- Anyone logging into the system instantly sees the real-time state of every slot across the faculty — no need to ask around or cross-reference separate sheets
- Lecturers avoid wasting time on slots that are already Reserved by another lecturer
- PLs see all pending requests in one place without chasing down email threads
- The database is the single source of truth — no conflicting versions of reality

Advantage 5 — Structured Approval Workflow with Audit Trail

- Google Sheets: PL manually writes "Y" in a cell, no automatic logging
 - Proposed system: FCFS dashboard with one-click approval + full timestamped audit trail
 - Benefit: complete accountability for all replacement decisions
-

1.4.2 Contributions (Stakeholders, SDG, Commercialisation)

Contribution to Lecturers

- Eliminates manual cross-referencing of schedules and venue PDFs
- Provides instant visibility of available slots with clear state indicators
- Reduces time spent on replacement administration

Contribution to Programme Leaders

- Pre-validated queue removes need for manual re-verification
- One-click approval with automatic audit trail
- Clear chronological overview of all pending requests

Contribution to Students

- View-only access to consolidated master timetable
- Real-time visibility of replacement status and schedule changes
- Ensures schedule integrity through OCC prevention of conflicts

Contribution to the Institution (FOCS Sabah)

- First system at TAR UMT Sabah to handle cross-cohort class replacement with OCC

- Establishes a technical precedent for concurrency-aware scheduling applications

Alignment with SDG 4 — Quality Education

- Reduces disruption to teaching schedules through rapid replacement processing
- Ensures continuity of learning activities when replacement is needed

Alignment with SDG 8 — Decent Work and Economic Growth

- Automates administrative coordination, freeing academic staff for teaching and mentoring
- Demonstrates integrity in educational resource management through OCC

Commercialisation and Showcase Potential

Proposed adoption pipeline:

- **Phase 1:** Deploy within FOCS TAR UMT Sabah (initial target — already scoped for this)
- **Phase 2:** Expand to all faculties at TAR UMT Sabah campus
- **Phase 3:** Adoption by TAR UMT main campus (KL) and relevant departments
- **Phase 4:** Roll out across all TAR UMT branches nationwide
- **Phase 5:** Adapt for other universities facing similar scheduling challenges

Showcase opportunities:

- Candidate for TAR UMT Smart Campus Initiative presentation
 - Technology showcase at faculty or institutional events
 - Reference implementation for OCC-based scheduling in academic environments
-

1.5 Project Plan

← covers Feasibility (5 marks)

1.5.1 Functional Modules ← scope as system decomposition (like your senior)

1.5.2 Milestones (Gantt table)

1.5.3 Software Development Model (Agile)

1.5.1 Functional Modules

Opening:

The system is decomposed into five functional modules. Each module encapsulates a distinct set of responsibilities, enabling modular development, testing, and

maintenance throughout the FYP1–FYP2 timeline.

Module 1 — Schedule and Timetable Management Module

- Stores and manages timetable data for lecturers, cohorts, and rooms
- Maintains master records for: 5 RSD3 cohorts, 5 staff members, 14 rooms
- Defines venue metadata: capacity, room type (Tutorial/Lecture Hall/Lab/Cisco Lab), allowed session types (L/T/P)
- Provides the data foundation for all other modules
- **Related objective:** 1.1.5 (Deployment)

Module 2 — Multi-Entity Matrix Intersection Module

- Core computational layer implementing the 4-vector set intersection algorithm
- Inputs: lecturer availability, 1+ cohort free schedules, room occupancy (from selected venue in dropdown), capacity filter
- Venue selection: dropdown above the timetable grid, defaulted to the class's **original venue**
- Changing the venue dropdown recalculates available slots in real-time
- Output: a weekly timetable grid (time × day) with color-coded cells:
 - **Green** — slot satisfies all 4 constraints (lecturer × cohort(s) × room × capacity → clickable)
 - **Red** — slot is occupied by an existing class
 - **Yellow** — slot has a pending request you submitted
 - **Grey** — slot is reserved by another lecturer's submission
 - **Blue** — Current Selection
- **Related objective:** 1.1.1 (Matrix Engine)

Module 3 — Optimistic Concurrency Control (OCC) Module

- Transactional validation layer that checks slot state at the precise millisecond of database write
- Handles: concurrent submission detection, rollback on conflict, user-facing alert generation
- Manages the slot state machine: Available ↔ Pending ↔ Reserved ↔ Occupied
- **Related objective:** 1.1.2 (OCC Layer)

Module 4 — FCFS Approval Dashboard Module

- Chronologically ordered queue of all pending replacement requests
- One-click approval (slot → Occupied) and rejection with reason (slot → Available)
- Full audit trail: PL identity, timestamp, action, slot ID, rejection reason

- **Related objective:** 1.1.3 (FCFS Dashboard)

Module 5 — Role-Based Access Control (RBAC) Module

- Authentication (username/password) and session management
- Three role tiers: Student (view-only), Lecturer (draft/submit), PL (hybrid)
- Authorization gates on all routes and actions
- **Related objective:** 1.1.4 (RBAC)

1.5.2 Milestones (Gantt Table)

Opening:

The development follows an Agile iterative approach across FYP1 (14 weeks) and FYP2 (7 weeks). High-risk features (schedule matching, OCC) are prioritised early.

Phase	Activity	Expected Outcome	Target Date
FYP1	Requirements Elicitation and Problem Formulation	Approved Form 2 Proposal	FYP1 Week 2
FYP1	Chapter 1: Introduction	Completed problem statement, objectives, scope	FYP1 Week 5
FYP1	System Requirements and Use Case Modelling	SRS, use case diagrams	FYP1 Week 6
FYP1	Chapter 2: Literature Review	APA-cited review of OCC, timetabling algorithms	FYP1 Week 8
FYP1	Database Schema Design	Normalised PostgreSQL schema, ERD	FYP1 Week 9
FYP1	Chapter 3: Methodology and Requirements	Functional + non-functional requirements	FYP1 Week 10
FYP1	Chapter 4: System Design	Architecture, UI mockups, OCC pseudocode	FYP1 Week 12
FYP1	Interim Prototype	Routing framework, seed data	FYP1 Week 13
FYP1	FYP1 Portfolio Submission	Chapters 1–4 report, viva presentation	FYP1 Week 14

Phase	Activity	Expected Outcome	Target Date
FYP2	Sprint 1: Matrix Intersection Engine	Functional core API for slot intersection	FYP2 Week 2
FYP2	Sprint 2: OCC + FCFS Dashboard	Conflict prevention, approval queue	FYP2 Week 4
FYP2	Sprint 3: UI Integration + RBAC	Complete interface, test report	FYP2 Week 5
FYP2	Final Thesis Compilation	Complete report (Chapters 5–7)	FYP2 Week 6
FYP2	Final Viva Defence	Presentation + live demo	FYP2 Week 7

1.5.3 Software Development Model — Agile (Iterative Approach)

Opening:

Due to the limited development period, the Agile iterative development method is adopted. This approach supports rapid development through short, repeated improvement cycles, allowing early testing of critical system components.

Why Agile fits this project:

- **Limited timeframe:** FYP2 has only 7 weeks — Agile's time-boxed iterations (sprints) fit naturally
- **High-risk features prioritised early:** The Matrix Intersection Engine and OCC layer are developed in Sprint 1 and Sprint 2 respectively, ensuring the most technically challenging components have the most testing time
- **Incremental delivery:** Each sprint produces a working, testable increment — Sprint 1 delivers the core API, Sprint 2 adds transactional integrity, Sprint 3 completes the UI
- **Adaptability:** Requirements can be refined based on supervisor feedback during progress meetings
- **Continuous validation:** Each increment can be demonstrated and validated before the next sprint begins

Sprint breakdown:

- **Sprint 1 (FYP2 W1–2):** Matrix Intersection Engine — database schema, timetable import, intersection algorithm API

- **Sprint 2 (FYP2 W3–4):** OCC Layer + FCFS Dashboard — transaction handling, approval queue, audit trail
- **Sprint 3 (FYP2 W5):** UI Integration + RBAC — Blade views, TailwindCSS styling, role gating, test cases
- **Sprint 4 (FYP2 W6–7):** Final thesis compilation, system documentation, viva preparation